

Concurrency for Scientific Computing in Python

AM215 - LECTURE 10

Today's Mission

Build efficient, scalable scientific workflows by choosing the right concurrency model.

We will cover:

1. **The "Why" of Concurrency:** CPU-bound vs. I/O-bound problems.
2. **Multiprocessing:** True parallelism for CPU-bound tasks.
3. **Threading & The GIL:** Overlapping I/O-bound tasks.
4. **Asyncio:** Massive concurrency for modern I/O.
5. **Choosing the Right Tool:** A practical decision framework.

Part 1: The Limitations of Synchronous Code

Synchronous Execution: One Step at a Time

By default, Python code runs sequentially. Each task must finish before the next one begins.

```
import time

def task_one():
    print("Starting task one...")
    time.sleep(1) # Simulate work
    print("Task one finished.")

def task_two():
    print("Starting task two...")
    time.sleep(1) # Simulate work
    print("Task two finished.")

task_one()
task_two()
# Total time: ~2 seconds
```

This one-at-a-time model is simple and predictable, but it becomes a major bottleneck when tasks involve significant work or waiting.

Bottleneck 1: CPU-Bound Work



■ Tasks limited by the speed of the processor. The CPU is the bottleneck.

- The program is actively performing computations, keeping the CPU busy.
- Examples:
 - Complex numerical calculations (e.g., matrix multiplication).
 - Intensive data processing and transformations.
 - Monte Carlo simulations.

| For these tasks, performance is determined by how fast the CPU can execute instructions. A synchronous program can only use one CPU core at a time.

Bottleneck 2: I/O-Bound Work

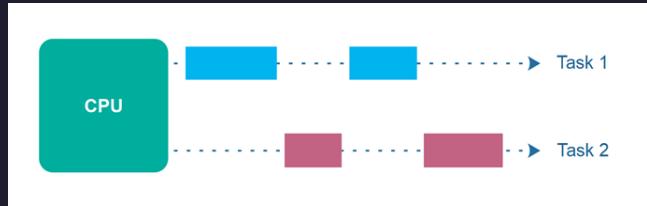


Tasks limited by waiting for slow external resources. The CPU is mostly idle.

- The program spends most of its time waiting for Input/Output (I/O) operations to complete.
- Examples:
 - Network requests (e.g., downloading a file, calling an API).
 - Reading or writing large files from a disk.
 - Querying a database.

For these tasks, performance is determined by the speed of the network or disk. A synchronous program wastes CPU time by sitting idle during these waits.

Concept: Concurrency

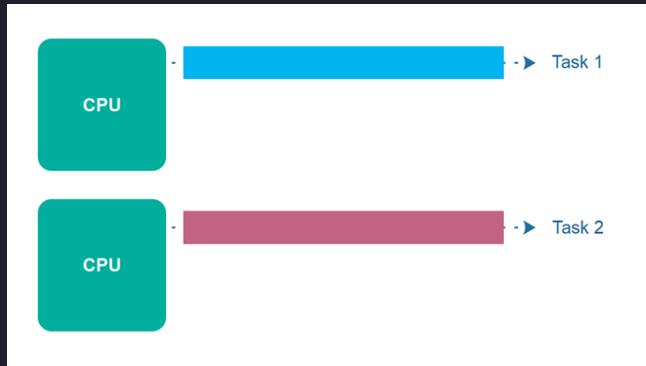


■ Structuring a program to manage multiple tasks at once.

- Tasks can be started, run, and completed in overlapping time periods.
- They are not necessarily running at the same instant.
- **Analogy:** A chef making a salad while also occasionally stirring a soup. One person is making progress on two tasks by switching between them.

■ Concurrency is about dealing with lots of things at once.

Concept: Parallelism



Executing multiple tasks at the exact same time.

- This is a stronger guarantee than concurrency. It requires hardware with multiple cores or processors.
- **Analogy:** Two chefs, each in their own kitchen, making their own dish. They are working simultaneously.

Parallelism is about doing lots of things at once.

Concurrency vs. Parallelism: A Summary



Concurrency is about structure; parallelism is about execution.

Concurrency

- **Goal:** Manage multiple tasks.
- **How:** Interleaving execution.
- **Requires:** A single core is sufficient.

Parallelism

- **Goal:** Execute multiple tasks simultaneously.
- **How:** Running on separate cores.
- **Requires:** Multiple hardware cores.

You can have concurrency without parallelism (e.g., on a single-core CPU), but you cannot have parallelism without concurrency.

The Process: An Isolated Program

■ A process is an instance of a computer program that is being executed.

- **Isolation:** Each process has its own private memory space. One process cannot directly access the memory of another.
- **Resources:** It has its own set of resources, such as memory, file handles, and network connections.
- **Cost:** Processes are "heavyweight" for the operating system to create and manage.

| Think of a process as a separate, self-contained workshop.

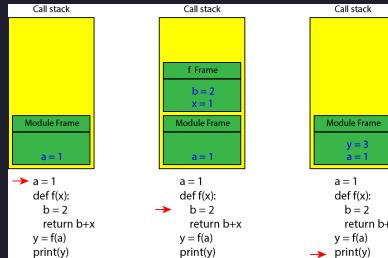
The Thread: A Flow Inside a Process

■ A thread is the smallest sequence of programmed instructions that can be managed independently by a scheduler.

- A single process can contain multiple threads.
- **Shared Memory:** All threads within a process share the same memory space. This makes communication between them fast, but also dangerous if not managed carefully.
- **Cost:** Threads are "lightweight" to create compared to processes.

Think of threads as multiple workers inside the same workshop. They can collaborate on the same project but must coordinate to avoid getting in each other's way.

Concept: The Thread's Stack

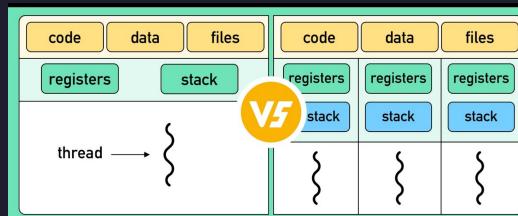


Each thread has its own call stack.

- **Call Stack:** A data structure that stores information about the active subroutines of a computer program. It's where local variables and function call information are kept.
- **Isolation within a Process:** While threads *share* the process's main memory (heap), each thread maintains its *own separate stack*.
- **Purpose:** This allows each thread to execute functions independently, keeping track of its own execution path and local variables without interfering with other threads.

Think of it as each worker in the workshop having their own personal workbench where they keep their current tools and instructions.

Process vs. Thread: Key Differences



■ The choice between them depends on memory sharing and the type of work.

Process

- **Memory:** Isolated.
- **Communication:** Slow (IPC).
- **Cost:** High.
- **Best for:** CPU-bound tasks where memory isolation is a benefit.

Thread

- **Memory:** Shared.
- **Communication:** Fast (shared variables).
- **Cost:** Low.
- **Best for:** I/O-bound tasks where overlapping waits is key.

Part 2: Multiprocessing for CPU-Bound Tasks

Motivator: 10 Million Draws Before Coffee

 Our first challenge: A classic CPU-bound scientific computing task.

- **Scenario:** We need to run a Monte Carlo simulation with 1,000 different parameter settings, where each setting requires 10,000 independent draws.
- **Problem:** This is a "Lots of Doing" task. Sequentially, it will take a long time as it runs on a single CPU core.

Question: Can the work for each parameter setting be done independently? If so, how can we leverage a multi-core CPU to speed this up?

Baseline: Sequential Monte Carlo

■ The straightforward approach uses a single core.

```
import time

def trial(n_samples):
    # Simulate n_samples and return a result
    # (This is a CPU-intensive calculation)
    ...

    # Run 1000 tasks sequentially
t0 = time.perf_counter()
results = [trial(10_000) for _ in range(1000)]
print(f"Sequential time: {time.perf_counter() - t0:.2f}s")
```

- This code is simple and correct, but inefficient.
- It cannot take advantage of modern multi-core processors.



A Naive Approach: Manual Parallelization

▀▀▀ What if we just run the script multiple times from the terminal?

```
# We could modify our script to take a command-line argument...
# python run_mc.py --chunk=1 &
# python run_mc.py --chunk=2 &
# ...
```

- This *is* a form of multiprocessing, but it's manual, clumsy, and hard to scale.
- We have to manually aggregate the results from different runs.
- This highlights the need for a tool that can manage this process for us from within a single Python script.

The Tool: concurrent.futures.ProcessPoolExecutor

■ Python's high-level interface for true parallel execution.

- It creates and manages a pool of separate Python **processes**.
- The operating system can schedule these processes to run on different CPU cores simultaneously.
- This is the standard way to achieve **parallelism** for CPU-bound tasks in Python.

```
from concurrent.futures import ProcessPoolExecutor  
  
tasks = [10_000] * 1000  
  
with ProcessPoolExecutor() as executor:  
    results = list(executor.map(trial, tasks))
```

- `with ProcessPoolExecutor():` ...: The `with` statement creates the process pool and ensures it is properly shut down when we're done, even if errors occur.
- `executor.map(trial, tasks):` This is the core of the operation. It "maps" the `trial` function onto each item in the `tasks` list, distributing the calls among the worker processes in the pool.

Live Run: Monte Carlo with Processes

Let's see the speedup from parallel execution.

```
# mp_pi.py
import time, random
from concurrent.futures import ProcessPoolExecutor

def hits(n):
    """Simulates n random draws to estimate pi."""
    r = random.random; c = 0
    for _ in range(n):
        x, y = r(), r(); c += (x*x + y*y <= 1.0)
    return c

if __name__ == "__main__":
    TASKS, N = 8, 1_000_000
    t0 = time.perf_counter()
    with ProcessPoolExecutor() as ex:
        total = sum(ex.map(hits, [N]*TASKS))
    print(f"pi≈{4*total/(TASKS*N):.5f}, time={time.perf_counter()-t0:.3f}s")
```

By distributing the `hits` function across multiple processes, we achieve a significant speedup that scales with the number of available CPU cores.

Cost 1: Startup Overhead

 Creating a new process is expensive for the operating system.

- Each worker process is a full, separate Python interpreter that must be launched and initialized.
- This cost is particularly high on macOS and Windows, which use the `spawn` start method by default. A `fork` (on Linux) is faster but has its own safety issues.
- **Impact:** Parallelizing very fast function calls can be slower than running them sequentially, as the startup cost dominates.

Cost 2: Data Serialization (Pickling)

 Data sent between processes must be converted into a byte stream.

- Processes don't share memory. To send an object (like a function argument or a return value) to another process, Python must **serialize** it using the `pickle` module.
- The receiving process must then **deserialize** it.
- **Impact:** This adds overhead that can be significant for large or complex Python objects (e.g., large lists, custom class instances).

Cost 3: Memory Duplication

 Each process has its own private memory space.

- When you pass a large data structure (like a big NumPy array or Pandas DataFrame) to a worker process, a copy of that data may be created in the worker's memory.
- If you have 8 worker processes, you could end up with 8 copies of your data in RAM.
- **Impact:** Can lead to very high memory consumption and even system crashes if not managed carefully.

Multiprocessing: Best Practices

▀▀▀ How to mitigate the costs and write robust parallel code.

1. Use Fewer, Larger Tasks

- The ideal task is long enough to make the startup and serialization costs negligible.
- Don't parallelize very short function calls; the overhead will dominate.

2. Guard Your Main Execution

- Always protect your main script's logic with `if __name__ == "__main__":`
- On `spawn` systems (macOS/Windows), worker processes re-import the main script. Without this guard, each worker would try to start its own process pool, leading to an infinite loop of process creation.

3. Manage Large Data Carefully

- Avoid passing large datasets as arguments if possible.
- For truly large arrays (e.g., NumPy), consider using `shared_memory` or memory-mapped files to allow processes to access the same data without copying it.

Part 3: Threading & The GIL

Motivator: 500 Tiny Files from the Cloud

■ Our second challenge: A classic I/O-bound task.

- **Scenario:** We need to download 500 small JSON files from a web API. Each download takes ~100ms due to network latency, but only ~2ms of CPU time to parse.
- **Problem:** This is a "Lots of Waiting" task. Sequentially, the CPU is idle >98% of the time, waiting for the network. Total time is dominated by the sum of all wait times.

| **Question:** Can we overlap the "waiting" periods? Can we start a new download while a previous one is in flight?

Baseline: Sequential File Fetching

▀▀▀ The straightforward approach waits for each download to complete.

```
import time, requests

URLS = ["https://example.com"] * 50

def fetch(url):
    r = requests.get(url, timeout=10)
    return r.content

t0 = time.perf_counter()
results = [fetch(u) for u in URLs]
print(f"Sequential time: {time.perf_counter() - t0:.2f}s")
```

- This code is simple, but its total runtime is the sum of all individual network latencies.
- The CPU sits idle while the operating system waits for network packets.

The Tool: concurrent.futures.ThreadPoolExecutor

Python's high-level interface for concurrent execution.

- It creates and manages a pool of **threads** within a single process.
- Threads are well-suited for I/O-bound tasks because they can overlap periods of waiting.
- While one thread is blocked (e.g., waiting for the network), the OS can switch to another thread to start a new task.

```
from concurrent.futures import ThreadPoolExecutor  
  
with ThreadPoolExecutor(max_workers=20) as executor:  
    results = list(executor.map(fetch, URLs))
```

- **max_workers=20**: Sets the size of the thread pool. For I/O-bound tasks, this can be larger than the number of CPU cores. Finding the optimal number often requires experimentation.
- **executor.map(fetch, URLs)**: Just like with **ProcessPoolExecutor**, this distributes the **fetch** calls across the threads in the pool.

This allows many I/O operations to be "in flight" at the same time, dramatically reducing the total wall-clock time.

Live Run: I/O Threads *Win* on Waiting

Let's see the speedup from overlapping I/O.

```
# io_threads_ok.py
import time, requests
from concurrent.futures import ThreadPoolExecutor

URLS = ["https://example.com"]*50

def fetch(u):
    r = requests.get(u, timeout=10)
    return len(r.content)

t0 = time.perf_counter()
with ThreadPoolExecutor(max_workers=20) as ex:
    sizes = list(ex.map(fetch, URLS))
print(f"bytes={sum(sizes)}, time={time.perf_counter()-t0:.3f}s")
```

By overlapping the network wait times, threads provide a massive performance boost for I/O-bound workloads. The total time is now closer to the duration of the longest few requests, not the sum of all of them.

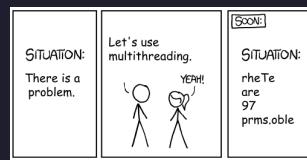
The Twist: What Happens with CPU-Bound Work?

We saw threads excel at I/O. Let's try them on our CPU-bound Monte Carlo task.

- **Hypothesis:** If we have 8 CPU cores, using 8 threads should give us a similar speedup to using 8 processes, right?



This leads us to one of the most important and misunderstood concepts in Python concurrency: The Global Interpreter Lock.



The Challenge of Shared State: Race Conditions

When multiple threads access and modify shared data, the final result can depend on the unpredictable order of execution.

```
import threading

counter = 0

def increment():
    global counter
    # This is not an atomic operation!
    # It involves: 1. read, 2. increment, 3. write.
    counter += 1

threads = [threading.Thread(target=increment) for _ in range(100_000)]
for t in threads: t.start()
for t in threads: t.join()

print(f"Final counter: {counter}") # Expected: 100000, Actual: ???
```

A **race condition** occurs because a thread can be interrupted between reading the value of `counter` and writing the new value back. This leads to lost updates and incorrect results.

Solution: Mutexes (Locks)

■ A mutex (Mutual Exclusion lock) is a synchronization primitive that ensures only one thread can execute a critical section of code at a time.

```
import threading

counter = 0
lock = threading.Lock()

def increment():
    global counter
    with lock:
        # The 'with' statement automatically acquires and releases the lock.
        # Only one thread can be inside this block at a time.
        counter += 1

    # ... (same thread setup as before) ...
    print(f"Final counter: {counter}") # Expected: 100000, Actual: 100000
```

By wrapping the critical section (`counter += 1`) in a lock, we make the operation **atomic** and prevent race conditions, ensuring correctness at the cost of some performance.



The Global Interpreter Lock (GIL)



■ A lock that protects CPython's internal memory management.

- **What it is:** A single, process-wide mutex that allows only one **thread** to execute Python bytecode at any given time, even on a multi-core processor.
- **Why it exists:** It simplifies CPython's memory management (e.g., reference counting) by acting as a global lock that prevents race conditions when multiple threads try to modify the same Python object simultaneously.

The GIL is specific to the CPython interpreter. Other implementations like Jython or IronPython do not have a GIL.

GIL Impact on CPU-Bound Python

■ The GIL prevents multiple threads from running Python code in parallel.

- The operating system can switch between threads on a multi-core machine, but only the thread that holds the GIL can actually execute Python bytecode.
- **Result:** No performance gain for pure Python computation. The overhead of thread management often makes it slightly *slower* than a simple sequential implementation.



GIL Impact on I/O-Bound Python



Threads can achieve true concurrency for I/O because they release the GIL.

- A **blocking I/O call** is an operation (like a network request or disk read) that causes the program to pause and wait for the external resource.
- In CPython, the standard library is designed so that when a thread makes a blocking I/O call, it **releases the GIL**.
- **Result:** While one thread is waiting for the network, the OS can switch to another thread, which can then acquire the GIL and run. This enables many I/O operations to be overlapped.

The GIL and Scientific Libraries (NumPy, SciPy)



C extensions can release the GIL, enabling parallelism for scientific code.

- Many scientific libraries like NumPy, SciPy, and Pandas are written in C or Fortran for performance.
- When you call a function like `np.dot()`, you are executing compiled C code, not Python bytecode.
- These libraries are often designed to **release the GIL** during long-running computations.
- **Result:** You can use a thread to run a heavy NumPy calculation on one core, and while it's running, another thread can acquire the GIL to do other work (like more Python processing or I/O). This is a powerful pattern for mixed workloads.

Live Run: When Threads Disappoint

Let's re-run our CPU-bound `hits` function with a `ThreadPoolExecutor`.

```
# cpu_threads_bad.py
import time
from concurrent.futures import ThreadPoolExecutor

def hits(n):
    """Simulates n random draws to estimate pi."""
    r = random.random; c = 0
    for _ in range(n):
        x, y = r(), r(); c += (x*x + y*y <= 1.0)
    return c

if __name__ == "__main__":
    TASKS, N = 8, 1_000_000
    t0 = time.perf_counter()
    with ThreadPoolExecutor(max_workers=TASKS) as ex:
        total = sum(ex.map(hits, [N]*TASKS))
    print(f"Threads (CPU-bound) time: {time.perf_counter()-t0:.3f}s")
```

Because `hits` is pure Python code, the GIL prevents parallel execution. The threads run one at a time, and performance is similar to (or worse than) the sequential version. This is the critical difference between multiprocessing and threading.

A Note on Testing Concurrent Code



Testing threaded code is notoriously difficult.

- **Non-Determinism:** Race conditions may only appear under specific, unpredictable timing conditions. A test suite might pass 1,000 times and then fail on the 1,001st run.
- **Heisenbugs:** Bugs that seem to disappear or change when you try to debug them (e.g., by adding print statements, which can alter the timing).
- **Strategies:**
 - **Stress Testing:** Running tests repeatedly in a loop to increase the chance of triggering a race condition.
 - **Fuzzing:** Using tools to inject random delays or inputs to explore unlikely execution paths.
 - **Careful Design:** The best approach is to minimize shared state and use well-tested synchronization primitives (like `queue.Queue`) to communicate between threads.

Part 4: Asyncio for Massive Concurrent I/O

The Problem with Threads at Scale

Threads are effective, but have limitations for massive concurrency.

- **Memory Footprint (Stack):** Each thread requires its own dedicated `call stack` (typically 1-8 MB). Creating thousands of threads can quickly consume gigabytes of RAM.
- **OS Overhead (Scheduling):** The operating system must manage and schedule each thread. With thousands of threads, the overhead of context switching and managing these resources can degrade performance.

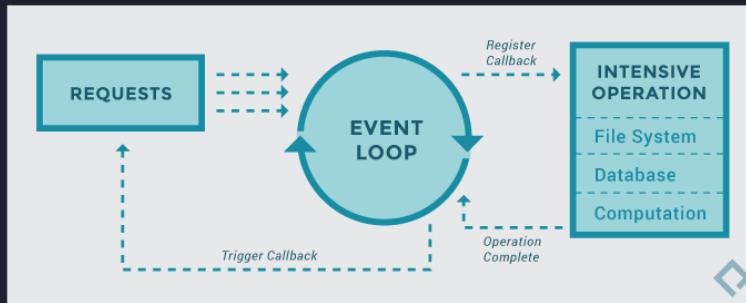
For applications needing to handle thousands of simultaneous network connections (like a modern web server or data streaming service), a more lightweight model is required.

Asyncio: The Event Loop & Cooperative Multitasking

- A single-threaded approach to concurrency.

Instead of relying on the OS to switch between threads, asyncio uses a single thread and an **event loop**.

- **Event Loop:** A coordinator that manages and distributes events or messages in a program. It keeps track of all running tasks.
- **Cooperative Multitasking:** Tasks explicitly **yield control** back to the event loop when they encounter a blocking I/O operation. The event loop can then run another task.



Asyncio Vocabulary: `async` and `await`

 New syntax for defining and controlling cooperative tasks.

`async def`

`await`

- Defines a function as a **coroutine**.
- A coroutine is a special function that can be paused and resumed. It does not run until it is scheduled on an event loop.

- The keyword used *inside* a coroutine to pause its execution and pass control back to the event loop.
- You can only `await` other "awaitable" objects (e.g., another coroutine, or an I/O operation from an `async`-compatible library).

A **Task** is a coroutine that has been scheduled to run on the event loop. `asyncio.run(main())` starts the event loop and runs the main coroutine.

Live Run: Asyncio Fan-Out

Let's demonstrate non-blocking waits with `asyncio.sleep`.

`asyncio.sleep(t)` is an `awaitable` that completes after `t` seconds without blocking the event loop, allowing other tasks to run.

```
# aio_fanout.py
import asyncio, time

async def ioish(i):
    await asyncio.sleep(0.1) # Simulate a non-blocking 100ms I/O wait
    return i

async def main(n=30):
    t0 = time.perf_counter()
    # asyncio.gather runs many coroutines concurrently
    await asyncio.gather(*[ioish(i) for i in range(n)])
    print(f"asyncio with {n} tasks in {time.perf_counter()-t0:.3f}s")

if __name__ == "__main__":
    asyncio.run(main())
```

The total time is just over 0.1s, not 3.0s. `asyncio` runs all 30 "waits" concurrently on a single thread, demonstrating its efficiency for I/O-bound workloads without the overhead of creating 30 separate threads.

Part 5: Choosing the Right Tool



Decision Table: A High-Level Guide

Match the tool to the bottleneck.

Problem Type	Multiprocessing	Threading	Asyncio
CPU-Bound (Numerical computation)	✓ Excellent (True Parallelism)	✗ No (due to GIL)	✗ No
I/O-Bound (Network, Disk)	⚠ Okay (High overhead)	✓ Good (OS-managed)	✓ Excellent (Lightweight)
Massive I/O (1000s of connections)	✗ No (Too heavyweight)	⚠ Okay (High memory/OS cost)	✓ Excellent (Designed for this)

Decision Table: Detailed Comparison

■ A deeper look at the trade-offs.

	Multiprocessing	Threading	Asyncio
Best For	CPU-bound tasks (Parallel computation)	I/O-bound tasks (e.g., 10-100 requests)	Massive I/O-bound tasks (e.g., 1000s of connections)
Key Cost	High startup & memory	GIL prevents parallelism	Code complexity (<code>async/await</code>)
Execution	Parallel (multiple cores)	Concurrent (single core)	Concurrent (single core)
Memory	Separate memory per process	Shared memory per process	Shared memory (single process)

What to Try First: A Hierarchy of Optimization

 Concurrency is a powerful tool, but it's not the first one you should reach for.

1. Vectorize Your Code

- Before anything else, use NumPy, SciPy, Pandas, or PyTorch to replace slow Python loops with fast, compiled C or Fortran code.
- This is often the single biggest performance gain.

2. Profile Your Code

- Use a profiler to find the *actual* bottleneck. Is your code CPU-bound or I/O-bound?
- Don't guess! Measure.

3. Choose a Concurrency Model

- Only after vectorizing and profiling should you choose a concurrency model based on the bottleneck you identified.

Key Takeaways

Designing for Scale and Reproducibility

- **CPU-Bound?** → **Multiprocessing**. Use `ProcessPoolExecutor` for true parallelism across multiple cores. Watch out for data serialization costs.
- **I/O-Bound?** → **Threading**. Use `ThreadPoolExecutor` to overlap "waiting" time. The GIL prevents it from speeding up CPU-bound Python code.
- **Massive I/O?** → **Asynio**. Use `async/await` for the most lightweight, scalable concurrency, but be prepared for increased code complexity.
- **Always Profile First**. Don't optimize prematurely. Identify your bottleneck, then choose the right tool for the job.

Resources & Further Reading

- **concurrent.futures Documentation**
 - docs.python.org/3/library/concurrent.futures.html (*https://docs.python.org/3/library/concurrent.futures.html*)
- **asyncio Documentation**
 - docs.python.org/3/library/asyncio.html (*https://docs.python.org/3/library/asyncio.html*)
- **Real Python Guides**
 - **Intro to Python Concurrency** / (*https://realpython.com/python-concurrency/*)
 - **AsyncIO** (*https://realpython.com/async-io-python/*)
 - **Intro to Python Threading** (*https://realpython.com/intro-to-python-threading/*)
 - **Thread Locks** (*https://realpython.com/python-thread-lock*)
 - **The Python GIL** (*https://realpython.com/python-gil/*)

Thank you!

